

LING570 Hw10: MaxEnt POS tagger

Due: 12/9

The example files are under /dropbox/21-22/570/hw10/examples/.

Q1 (55 points): Create a MaxEnt POS tagger, **maxent_tagger.sh**.

- The command line is: `maxent_tagger.sh train_file test_file rare_thres feat_thres output_dir`
- The `train_file` and `test_file` have the format (e.g., **test.word_pos**):
 $w_1/t_1 \ w_2/t_2 \ \dots \ w_n/t_n$
- `rare_thres` is an integer: any words (in the `train_file` and `test_file`) that appear LESS THAN `raw_thres` times in the **train_file** are treated as *rare words*, and features such as `pref=xx` and `suf=xx` should be used for rare words (see Table 1 in (Ratnaparkhi, 1996)).
- `feat_thres` is an integer: All the w_i features (i.e., `CurrentWord=xx` features), regardless of their frequency, should be kept. For all OTHER types of features, if a feature appears LESS THAN `feat_thres` in the **train_file**, that feature should be removed from the feature vectors.
- `output_dir` is a directory that stores the output files from the tagger. Your script should create the following files and store them under `output_dir`:
 - `train_voc` (e.g., **ex_train_voc**): the vocabulary that includes all the words appearing in `train_file`.
 - * The file has the format “word freq” where freq is the frequency of the word in the training data.
 - * The lines should be sorted by freq in descending order.
 - * For words with the same frequency, no need to sort the lines. We will sort your file before comparing it with the gold standard.
 - `init_feats` (e.g., **ex_init_feats**): features that occur in the `train_file`.
 - * It has the format “featName freq”.
 - * The lines are sorted by the frequency of the feature in the `train_file` in descending order.
 - * For features with the same frequency, no need to sort the lines.
 - `kept_feats` (e.g., **ex_kept_feats**): This is a subset of `init_feats`, and it includes the features that are kept after applying `feat_thres`. The file has the same format and order as in `init_feats`.
 - `final_train.vectors.txt` (e.g., **ex_final_train.vectors.txt**): the feat vectors for the `train_file` in the Mallet text format. Only features in `kept_feats` should be kept in this file.
 - `final_test.vectors.txt`: the feat vectors for the `test_file` in the Mallet text format.
 - * The format is the same as `final_train.vectors.txt`.
 - * It would be preferred that this file contains only features that appear in `kept_feats`.
 - `final_train.vectors`: the binary format of the vectors in `final_train.vectors.txt`.
 - `me_model`: the MaxEnt model (in binary format) which is produced by the MaxEnt trainer.

- `me_model.stdout` and `me_model.stderr`: the stdout (standard out) and stderr (standard error) produced by the MaxEnt trainer are redirected and saved to those files by running command such as “`mallet train-classifier --trainer MaxEnt --input final_train.vectors --output-classifier me_model > me_model.stdout 2 > me_model.stderr`”. The training accuracy is displayed at the end of `me_model.stdout`.
- `sys_out`: the system output file when running the MaxEnt classifier with command such as “`mallet classify-file --input final_test.vectors.txt --classifier me_model --output sys_out`”.

Your script **maxent_tagger.sh** should do the following:

1. Create feature vectors for the training data and the test data. The vector files should be called **final_train.vectors.txt** and **final_test.vectors.txt**.
2. Run **mallet import-file** to convert the training vectors into binary format, and the binary file is called **final_train.vectors**.
3. Run **mallet train-classifier** to create a MaxEnt model **me_model** using **final_train.vectors**
4. Run **mallet classify-file** to get the result on the test data **final_test.vectors.txt**.
5. Calculate the test accuracy

For step 2-4, you should use Mallet commands. For Step 5, if you don’t want to write code for it, you can use the `vectors2classify` command, which covers step 3-5. In that case, you need to convert **final_test.vectors.txt** to the binary format first.

For the first step, you need to write some code. Features are defined in Table 1 in (Ratnaparkhi, 1996). The following is one way for implementing this step:

1. create `train_voc` from the `train_file`, and use the word frequency in `train_voc` and `rare_thres` to determine whether a word should be treated as a *rare word*. The feature vectors for rare words and non-rare words are different.
2. Form feature vectors for the words in `train_file`, and store the features and frequencies in the **training data** in `init_feats`.
3. Create `kept_feats` by using `feat_thres` to filter out low frequency features in `init_feats`. Note that w_i features are NOT subject to filtering with `feat_thres` and every w_i feature in `init_feats` should be kept in `kept_feats`.
4. Go through the feature vector file for `train_file` and remove all the features that are not in `kept_feats`.
5. Create feature vectors for `test_file`, and use only the features in `kept_feats`. If a word in the `test_file` appears LESS THAN `rare_thres` times (or does not appear at all) in the `training_file`, the word should be treated as a *rare word* even if it appears many times in the `test_file`.
6. For the feature vector files, replace all the occurrences of “,” with “**comma**” as Mallet treats “,” as a separator.

Handling the beginning and the end of a sentence:

- When you create feature vector files, you need to do that only for words appearing in the sentences. No need to create feature vectors for the sentence beginning marker “</s>” and for the sentence ending marker “<s>”.
- For the first word in the sentence, the previous tag will be BOS, the previous word will be “</s>”. For the sake of simplicity, you can just use prevW=BOS for previous word feature, instead of prevW=</s>. Mallet will treat the two feature names the same way.
- Similarly, for the last word in the sentence, the next tag is EOS, and the next word is “<s>” (or you can call it EOS).

Q2 (20 points): Run maxent_tagger.sh with wsj_sec0.word_pos as train_file, test.word_pos as test_file, and the thresholds as specified in Table 1:

- *training accuracy* is the accuracy of the tagger on the train_file
- *test accuracy* is the accuracy of the tagger on the test_file
- *# of feats* is the number of features in the train_file before applying feat_thres
- *# of kept feats* is the number of features in the train_file after applying feat_thres
- *running time* is the CPU time (in minutes) of running maxent_tagger.sh.

Table 1: Tagging accuracy with different thresholds

Expt id	rare thres	feat thres	training accuracy	test accuracy	# of feats	# of kept feats	running time
1_1	1	1					
1_3	1	3					
2_3	2	3					
3_5	3	5					
5_10	5	10					

Please do the following:

- Fill out Table 1.
- What conclusion can you draw from Table 1?
- Save the output files of maxent_tagger.sh to res.id/, where id is the experiment id in the first column (e.g., the files for the first experiment will be stored under res_1_1). Submit only the subdirs for the first row and the last row (i.e., res_1_1 and res_5_10).

Submission: Your submission should include the following:

1. readme.[txt|pdf] includes Table 1 and your answer to Q2.
2. hw.tar.gz that includes maxent_tagger.sh and res_1_1/ and res_5_10/ created in Q2 (see the complete file list in submit-file-list).